

## Movie Analysis

Our goal is to identify:

- top genre combinations
- top producers
- directors
- actors
- actresses

We will focus on two metrics:

- average Return on Investment
- average profit

```
import sqlite3
import numpy as np
import pandas as pd
from fuzzywuzzy import fuzz
from fuzzywuzzy import process
import recordlinkage
from recordlinkage.preprocessing import clean
```

## IMDb database

Lets connect to the database and explore the data

```
path = "zippedData/im.db"
conn = sqlite3.connect(path)
cursor = conn.cursor()
```

# to see all the tables in the database.

```
imdb_df = pd.read_sql(
    """
    SELECT *
    FROM sqlite_master
    """
    , conn
)
```

```
imdb_df[imdb_df['type'] == 'table']
```

|   | type  | name          | tbl_name      | rootpage | sql                                                |
|---|-------|---------------|---------------|----------|----------------------------------------------------|
| 0 | table | movie_basics  | movie_basics  | 2        | CREATE TABLE "movie_basics" (\n"movie_id" TEXT...  |
| 1 | table | directors     | directors     | 3        | CREATE TABLE "directors" (\n"movie_id" TEXT,\n...  |
| 2 | table | known_for     | known_for     | 4        | CREATE TABLE "known_for" (\n"person_id" TEXT,\n... |
| 3 | table | movie_akas    | movie_akas    | 5        | CREATE TABLE "movie_akas" (\n"movie_id" TEXT,\n... |
| 4 | table | movie_ratings | movie_ratings | 6        | CREATE TABLE "movie_ratings" (\n"movie_id" TEX...  |
| 5 | table | persons       | persons       | 7        | CREATE TABLE "persons" (\n"person_id" TEXT,\n ...  |
| 6 | table | principals    | principals    | 8        | CREATE TABLE "principals" (\n"movie_id" TEXT,\n... |
| 7 | table | writers       | writers       | 9        | CREATE TABLE "writers" (\n"movie_id" TEXT,\n ...   |

```
query1 = """ SELECT * FROM movie_basics ORDER BY -start_year LIMIT 10"""
pd.read_sql(query1, conn)
```

|   | movie_id   | primary_title                             | original_title                            | start_year | runtime_minutes | genres                   |
|---|------------|-------------------------------------------|-------------------------------------------|------------|-----------------|--------------------------|
| 0 | tt5174640  | 100 Years                                 | 100 Years                                 | 2115       | None            | Drama                    |
| 1 | tt5637536  | Avatar 5                                  | Avatar 5                                  | 2027       | None            | Action,Adventure,Fantasy |
| 2 | tt10300398 | Untitled Star Wars Film                   | Untitled Star Wars Film                   | 2026       | None            | Fantasy                  |
| 3 | tt3095356  | Avatar 4                                  | Avatar 4                                  | 2025       | None            | Action,Adventure,Fantasy |
| 4 | tt10300396 | Untitled Star Wars Film                   | Untitled Star Wars Film                   | 2024       | None            | None                     |
| 5 | tt6149054  | Fantastic Beasts and Where to Find Them 5 | Fantastic Beasts and Where to Find Them 5 | 2024       | None            | Adventure,Family,Fantasy |
| 6 | tt10255736 | Untitled Marvel Project                   | Untitled Marvel Project                   | 2023       | None            | Action                   |
| 7 | tt10298848 | Untitled Disney Live-Action Project       | Untitled Disney Live-Action Project       | 2023       | None            | None                     |
| 8 | tt1757678  | Avatar 3                                  | Avatar 3                                  | 2023       | None            | Action,Adventure,Drama   |
| 9 | tt6258542  | Wraith of the Umbra and Eidolon II        | Wraith of the Umbra and Eidolon II        | 2023       | None            | Adventure,Drama,Fantasy  |

```
# get column info
cursor.execute("PRAGMA table_info(movie_basics)")
cursor.fetchall()
```

```
[(0, 'movie_id', 'TEXT', 0, None, 0),
 (1, 'primary_title', 'TEXT', 0, None, 0),
 (2, 'original_title', 'TEXT', 0, None, 0),
 (3, 'start_year', 'INTEGER', 0, None, 0),
 (4, 'runtime_minutes', 'REAL', 0, None, 0),
 (5, 'genres', 'TEXT', 0, None, 0)]
```

```
query4 = """ SELECT COUNT(*) FROM movie_basics"""
pd.read_sql(query4, conn)
```

|   | COUNT(*) |
|---|----------|
| 0 | 146144   |

```
query2 = """ SELECT * FROM movie_ratings LIMIT 5"""
pd.read_sql(query2, conn)
```

|   | movie_id   | averagerating | numvotes |
|---|------------|---------------|----------|
| 0 | tt10356526 | 8.3           | 31       |
| 1 | tt10384606 | 8.9           | 559      |
| 2 | tt1042974  | 6.4           | 20       |
| 3 | tt1043726  | 4.2           | 50352    |
| 4 | tt1060240  | 6.5           | 21       |

```
query11 = """ SELECT * FROM directors LIMIT 5"""
pd.read_sql(query11, conn)
```

|   | movie_id  | person_id |
|---|-----------|-----------|
| 0 | tt0285252 | nm0899854 |
| 1 | tt0462036 | nm1940585 |
| 2 | tt0835418 | nm0151540 |
| 3 | tt0835418 | nm0151540 |
| 4 | tt0878654 | nm0089502 |

Let's combine the tables to include movie data, ratings, as well as the producers, directors, and actors who worked on the movie

```
query201 = """
WITH ranked_directors AS (
SELECT
    d.movie_id,
    p.primary_name AS director_name,
    ROW_NUMBER() OVER (PARTITION BY d.movie_id ORDER BY d.person_id) AS director_rank
```

```

FROM directors d
JOIN persons p ON d.person_id = p.person_id
GROUP BY d.movie_id, p.primary_name, d.person_id
),

ranked_principals AS (
  SELECT
    p.movie_id,
    per.primary_name AS person_name,
    p.category,
    ROW_NUMBER() OVER (PARTITION BY p.movie_id, p.category ORDER BY p.person_id) AS person_rank
  FROM principals p
  JOIN persons per ON p.person_id = per.person_id
  WHERE p.category IN ('actor', 'actress', 'producer')
)

SELECT
  mr.movie_id AS imdb_movie_id,
  mb.primary_title,
  mr.averagerating AS average_rating,
  mr.numvotes,
  mb.start_year,
  mb.runtime_minutes,
  mb.genres,
  MAX(CASE WHEN rd.director_rank = 1 THEN rd.director_name END) AS director1,
  MAX(CASE WHEN rd.director_rank = 2 THEN rd.director_name END) AS director2,
  MAX(CASE WHEN rp.category = 'actress' AND rp.person_rank = 1 THEN rp.person_name END) AS actress1,
  MAX(CASE WHEN rp.category = 'actress' AND rp.person_rank = 2 THEN rp.person_name END) AS actress2,
  MAX(CASE WHEN rp.category = 'actress' AND rp.person_rank = 3 THEN rp.person_name END) AS actress3,
  MAX(CASE WHEN rp.category = 'actor' AND rp.person_rank = 1 THEN rp.person_name END) AS actor1,
  MAX(CASE WHEN rp.category = 'actor' AND rp.person_rank = 2 THEN rp.person_name END) AS actor2,
  MAX(CASE WHEN rp.category = 'actor' AND rp.person_rank = 3 THEN rp.person_name END) AS actor3,
  MAX(CASE WHEN rp.category = 'actor' AND rp.person_rank = 4 THEN rp.person_name END) AS actor4,
  MAX(CASE WHEN rp.category = 'producer' AND rp.person_rank = 1 THEN rp.person_name END) AS producer1,
  MAX(CASE WHEN rp.category = 'producer' AND rp.person_rank = 2 THEN rp.person_name END) AS producer2
FROM movie_ratings mr
JOIN movie_basics mb ON mr.movie_id = mb.movie_id
LEFT JOIN ranked_directors rd ON mr.movie_id = rd.movie_id
LEFT JOIN ranked_principals rp ON mr.movie_id = rp.movie_id
WHERE mr.numvotes >= 30
GROUP BY
  mr.movie_id,
  mb.primary_title,
  mr.averagerating,
  mr.numvotes,
  mb.start_year,
  mb.runtime_minutes,
  mb.genres
ORDER BY -mr.numvotes

"""
pd.read_sql(query201, conn)

```

|       | imdb_movie_id | primary_title                           | average_rating | numvotes | start_year | runtime_minutes | genres                  | director1           | dir |
|-------|---------------|-----------------------------------------|----------------|----------|------------|-----------------|-------------------------|---------------------|-----|
| 0     | tt1375666     | Inception                               | 8.8            | 1841066  | 2010       | 148.0           | Action,Adventure,Sci-Fi | Christopher Nolan   |     |
| 1     | tt1345836     | The Dark Knight Rises                   | 8.4            | 1387769  | 2012       | 164.0           | Action,Thriller         | Christopher Nolan   |     |
| 2     | tt0816692     | Interstellar                            | 8.6            | 1299334  | 2014       | 169.0           | Adventure,Drama,Sci-Fi  | Christopher Nolan   |     |
| 3     | tt1853728     | Django Unchained                        | 8.4            | 1211405  | 2012       | 165.0           | Drama,Western           | Quentin Tarantino   |     |
| 4     | tt0848228     | The Avengers                            | 8.1            | 1183655  | 2012       | 143.0           | Action,Adventure,Sci-Fi | Joss Whedon         |     |
| ...   | ...           | ...                                     | ...            | ...      | ...        | ...             | ...                     | ...                 | ... |
| 43735 | tt9378760     | Sarah Millican: Control Enthusiast Live | 6.9            | 30       | 2018       | 82.0            | Comedy                  | Brian Klein         |     |
| 43736 | tt9442146     | Hüddam 2                                | 5.3            | 30       | 2019       | 92.0            | Drama,Horror,Thriller   | Utku Uçar           |     |
| 43737 | tt9598566     | Ave Maria                               | 7.3            | 30       | 2018       | 74.0            | Drama                   | Vipin Radhakrishnan |     |
| 43738 | tt9613316     | Frances Ferguson                        | 6.8            | 30       | 2019       | 74.0            | Comedy                  | Bob Byington        |     |
| 43739 | tt9647980     | Patria                                  | 7.5            | 30       | 2019       | 89.0            | Documentary             | Matías Gueilburt    |     |

43740 rows x 18 columns

```
imdb_df = pd.read_sql(query201, conn)
imdb_df.head()
```

|   | imdb_movie_id | primary_title         | average_rating | numvotes | start_year | runtime_minutes | genres                  | director1         | director2 |
|---|---------------|-----------------------|----------------|----------|------------|-----------------|-------------------------|-------------------|-----------|
| 0 | tt1375666     | Inception             | 8.8            | 1841066  | 2010       | 148.0           | Action,Adventure,Sci-Fi | Christopher Nolan | None      |
| 1 | tt1345836     | The Dark Knight Rises | 8.4            | 1387769  | 2012       | 164.0           | Action,Thriller         | Christopher Nolan | None      |
| 2 | tt0816692     | Interstellar          | 8.6            | 1299334  | 2014       | 169.0           | Adventure,Drama,Sci-Fi  | Christopher Nolan | None      |
| 3 | tt1853728     | Django Unchained      | 8.4            | 1211405  | 2012       | 165.0           | Drama,Western           | Quentin Tarantino | None      |
| 4 | tt0848228     | The Avengers          | 8.1            | 1183655  | 2012       | 143.0           | Action,Adventure,Sci-Fi | Joss Whedon       | None      |

```
imdb_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 43740 entries, 0 to 43739
Data columns (total 18 columns):
#   Column              Non-Null Count  Dtype
---  -
0   imdb_movie_id       43740 non-null  object
1   primary_title       43740 non-null  object
2   average_rating      43740 non-null  float64
3   numvotes            43740 non-null  int64
4   start_year          43740 non-null  int64
5   runtime_minutes     41466 non-null  float64
6   genres              43619 non-null  object
7   director1           43528 non-null  object
8   director2           4352 non-null   object
9   actress1            33042 non-null  object
10  actress2             19261 non-null  object
11  actress3             6159 non-null   object
12  actor1              37258 non-null  object
13  actor2              32216 non-null  object
14  actor3              19403 non-null  object
```

```

15 actor4          6332 non-null object
16 producer1      31705 non-null object
17 producer2      16412 non-null object
dtypes: float64(2), int64(2), object(14)
memory usage: 6.0+ MB

```

## ✓ Movie Budgets dataset

Let's explore the dataset containing movie budgets and revenue data

```

# Read the CSV file
budgets_df = pd.read_csv('zippedData/tn.movie_budgets.csv.gz')

# Display the first few rows of the DataFrame
print(budgets_df.tail(20))

```

```

↗ id release_date movie production_budget \
5762 63 Apr 11, 1997 Pink Flamingos $12,000
5763 64 Apr 28, 2006 Grip: A Criminal's Story $12,000
5764 65 Dec 31, 2007 Tin Can Man $12,000
5765 66 Mar 9, 2001 Dayereh $10,000
5766 67 Apr 28, 2006 Clean $10,000
5767 68 Jul 6, 2001 Cure $10,000
5768 69 May 28, 2004 On the Downlow $10,000
5769 70 Apr 1, 1996 Bang $10,000
5770 71 Aug 14, 2008 The Rise and Fall of Miss Thang $10,000
5771 72 May 19, 2015 Family Motocross $10,000
5772 73 Jan 13, 2012 Newlyweds $9,000
5773 74 Feb 26, 1993 El Mariachi $7,000
5774 75 Oct 8, 2004 Primer $7,000
5775 76 May 26, 2006 Cavite $7,000
5776 77 Dec 31, 2004 The Mongol King $7,000
5777 78 Dec 31, 2018 Red 11 $7,000
5778 79 Apr 2, 1999 Following $6,000
5779 80 Jul 13, 2005 Return to the Land of Wonders $5,000
5780 81 Sep 29, 2015 A Plague So Pleasant $1,400
5781 82 Aug 5, 2005 My Date With Drew $1,100

```

```

domestic_gross worldwide_gross
5762 $413,802 $413,802
5763 $1,336 $1,336
5764 $0 $0
5765 $673,780 $673,780
5766 $138,711 $138,711
5767 $94,596 $94,596
5768 $1,987 $1,987
5769 $527 $527
5770 $401 $401
5771 $0 $0
5772 $4,584 $4,584
5773 $2,040,920 $2,041,928
5774 $424,760 $841,926
5775 $70,071 $71,644
5776 $900 $900
5777 $0 $0
5778 $48,482 $240,495
5779 $1,338 $1,338
5780 $0 $0
5781 $181,041 $181,041

```

```
budgets_df.info()
```

```

↗ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 5782 entries, 0 to 5781
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     5782 non-null   int64
1   release_date          5782 non-null   object
2   movie                 5782 non-null   object
3   production_budget     5782 non-null   object
4   domestic_gross        5782 non-null   object
5   worldwide_gross       5782 non-null   object
dtypes: int64(1), object(5)
memory usage: 271.2+ KB

```

```
# convert release_dat column to datetime
budgets_df['release_date'] = pd.to_datetime(budgets_df['release_date'], format='%b %d, %Y')

# add a new column with just the year
budgets_df['release_year'] = budgets_df['release_date'].dt.year
budgets_df.head()
```

|   | id | release_date | movie                                       | production_budget | domestic_gross | worldwide_gross | release_year |
|---|----|--------------|---------------------------------------------|-------------------|----------------|-----------------|--------------|
| 0 | 1  | 2009-12-18   | Avatar                                      | \$425,000,000     | \$760,507,625  | \$2,776,345,279 | 2009         |
| 1 | 2  | 2011-05-20   | Pirates of the Caribbean: On Stranger Tides | \$410,600,000     | \$241,063,875  | \$1,045,663,875 | 2011         |
| 2 | 3  | 2019-06-07   | Dark Phoenix                                | \$350,000,000     | \$42,762,350   | \$149,762,350   | 2019         |
| 3 | 4  | 2015-05-01   | Avengers: Age of Ultron                     | \$330,600,000     | \$459,005,868  | \$1,403,013,963 | 2015         |
| 4 | 5  | 2017-12-15   | Star Wars Ep. VIII: The Last Jedi           | \$317,000,000     | \$620,181,382  | \$1,316,721,747 | 2017         |

```
budgets_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5782 entries, 0 to 5781
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    5782 non-null  int64
1   release_date          5782 non-null  datetime64[ns]
2   movie                 5782 non-null  object
3   production_budget     5782 non-null  object
4   domestic_gross        5782 non-null  object
5   worldwide_gross       5782 non-null  object
6   release_year          5782 non-null  int32
dtypes: datetime64[ns](1), int32(1), int64(1), object(4)
memory usage: 293.7+ KB
```

```
imdb_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 43740 entries, 0 to 43739
Data columns (total 18 columns):
#   Column                Non-Null Count  Dtype
---  -
0   imdb_movie_id         43740 non-null  object
1   primary_title         43740 non-null  object
2   average_rating        43740 non-null  float64
3   numvotes              43740 non-null  int64
4   start_year            43740 non-null  int64
5   runtime_minutes       41466 non-null  float64
6   genres                43619 non-null  object
7   director1             43528 non-null  object
8   director2             4352 non-null   object
9   actress1              33042 non-null  object
10  actress2              19261 non-null  object
11  actress3              6159 non-null   object
12  actor1                37258 non-null  object
13  actor2                32216 non-null  object
14  actor3                19403 non-null  object
15  actor4                6332 non-null   object
16  producer1             31705 non-null  object
17  producer2             16412 non-null  object
dtypes: float64(2), int64(2), object(14)
memory usage: 6.0+ MB
```

```
budgets_df['release_year'] = budgets_df['release_year'].astype('int64')
```

```
# Check for missing values in each column
budgets_df.isnull().sum()
```

```
id                0
release_date      0
movie             0
production_budget 0
domestic_gross    0
worldwide_gross   0
release_year      0
dtype: int64
```

```
imdb_df.isnull().sum()
```

```

imdb_movie_id      0
primary_title      0
average_rating     0
numvotes           0
start_year         0
runtime_minutes    2274
genres             121
director1          212
director2         39388
actress1          10698
actress2          24479
actress3          37581
actor1            6482
actor2           11524
actor3           24337
actor4           37408
producer1         12035
producer2         27328
dtype: int64

```

Start coding or [generate](#) with AI.

## ✓ Combining IMDb and Movie Budgets datasets

Normalize titles by converting to lowercase and stripping spaces

```

budgets_df['movie'] = budgets_df['movie'].str.lower().str.strip()
imdb_df['primary_title'] = imdb_df['primary_title'].str.lower().str.strip()

```

Years have to match exactly but not movie titles

```

# Create an indexer and define the comparison criteria
indexer = recordlinkage.Index()
indexer.block(left_on='release_year', right_on='start_year')
candidate_links = indexer.index(budgets_df, imdb_df)

compare = recordlinkage.Compare()
compare.exact('release_year', 'start_year', label='year')
compare.string('movie', 'primary_title', method='jarowinkler', threshold=0.95, label='title')
features = compare.compute(candidate_links, budgets_df, imdb_df)

# Filter matches based on a threshold
matches = features[features.sum(axis=1) > 1.5]

```

Now Let's merge the data frames based on the matches

```

budgets_df['index'] = budgets_df.index
imdb_df['index'] = imdb_df.index

matches.reset_index(inplace=True)
merged_df = pd.merge(budgets_df, matches, left_on='index', right_on='level_0')
merged_df = pd.merge(merged_df, imdb_df, left_on='level_1', right_on='index')

# Drop unnecessary columns and duplicates
merged_df.drop(columns=['index_x', 'index_y', 'level_0', 'level_1'], inplace=True)
merged_df.drop_duplicates(inplace=True)
merged_df

```



|      | id  | release_date | movie                                       | production_budget | domestic_gross | worldwide_gross | release_year | year | title | imdb_movie_id |
|------|-----|--------------|---------------------------------------------|-------------------|----------------|-----------------|--------------|------|-------|---------------|
| 0    | 2   | 2011-05-20   | pirates of the caribbean: on stranger tides | \$410,600,000     | \$241,063,875  | \$1,045,663,875 | 2011         | 1    | 1.0   | tt1298650     |
| 1    | 3   | 2019-06-07   | dark phoenix                                | \$350,000,000     | \$42,762,350   | \$149,762,350   | 2019         | 1    | 1.0   | tt6565702     |
| 2    | 4   | 2015-05-01   | avengers: age of ultron                     | \$330,600,000     | \$459,005,868  | \$1,403,013,963 | 2015         | 1    | 1.0   | tt2395427     |
| 3    | 7   | 2018-04-27   | avengers: infinity war                      | \$300,000,000     | \$678,815,482  | \$2,048,134,200 | 2018         | 1    | 1.0   | tt4154756     |
| 4    | 9   | 2017-11-17   | justice league                              | \$300,000,000     | \$229,024,295  | \$655,945,209   | 2017         | 1    | 1.0   | tt0974015     |
| ...  | ... | ...          | ...                                         | ...               | ...            | ...             | ...          | ...  | ...   | ...           |
| 1595 | 35  | 2013-10-25   | her cry: la llorona investigation           | \$35,000          | \$0            | \$0             | 2013         | 1    | 1.0   | tt2469216     |
| 1596 | 49  | 2015-09-01   | exeter                                      | \$25,000          | \$0            | \$489,792       | 2015         | 1    | 1.0   | tt1945044     |
| 1597 | 52  | 2015-12-01   | dutch kills                                 | \$25,000          | \$0            | \$0             | 2015         | 1    | 1.0   | tt2759066     |
| 1598 | 59  | 2011-11-25   | the ridges                                  | \$17,300          | \$0            | \$0             | 2011         | 1    | 1.0   | tt1781935     |
| 1599 | 62  | 2014-12-31   | stories of our lives                        | \$15,000          | \$0            | \$0             | 2014         | 1    | 1.0   | tt3973612     |

1600 rows x 27 columns

```
# rearrange columns
columns = list(merged_df.columns)

# Move 'primary_title' next to 'movie'
columns.insert(columns.index('movie') + 1, columns.pop(columns.index('primary_title')))
merged_df = merged_df[columns]
merged_df
```





|      | id  | release_date | movie                                       | primary_title                               | production_budget | domestic_gross | worldwide_gross | release_year | year | title |
|------|-----|--------------|---------------------------------------------|---------------------------------------------|-------------------|----------------|-----------------|--------------|------|-------|
| 0    | 2   | 2011-05-20   | pirates of the caribbean: on stranger tides | pirates of the caribbean: on stranger tides | \$410,600,000     | \$241,063,875  | \$1,045,663,875 | 2011         | 1    | 1.0   |
| 1    | 3   | 2019-06-07   | dark phoenix                                | dark phoenix                                | \$350,000,000     | \$42,762,350   | \$149,762,350   | 2019         | 1    | 1.0   |
| 2    | 4   | 2015-05-01   | avengers: age of ultron                     | avengers: age of ultron                     | \$330,600,000     | \$459,005,868  | \$1,403,013,963 | 2015         | 1    | 1.0   |
| 3    | 7   | 2018-04-27   | avengers: infinity war                      | avengers: infinity war                      | \$300,000,000     | \$678,815,482  | \$2,048,134,200 | 2018         | 1    | 1.0   |
| 4    | 9   | 2017-11-17   | justice league                              | justice league                              | \$300,000,000     | \$229,024,295  | \$655,945,209   | 2017         | 1    | 1.0   |
| ...  | ... | ...          | ...                                         | ...                                         | ...               | ...            | ...             | ...          | ...  | ...   |
| 1595 | 35  | 2013-10-25   | her cry: la llorona investigation           | her cry: la llorona investigation           | \$35,000          | \$0            | \$0             | 2013         | 1    | 1.0   |
| 1596 | 49  | 2015-09-01   | exeter                                      | exeter                                      | \$25,000          | \$0            | \$489,792       | 2015         | 1    | 1.0   |
| 1597 | 52  | 2015-12-01   | dutch kills                                 | dutch kills                                 | \$25,000          | \$0            | \$0             | 2015         | 1    | 1.0   |
| 1598 | 59  | 2011-11-25   | the ridges                                  | the ridges                                  | \$17,300          | \$0            | \$0             | 2011         | 1    | 1.0   |
| 1599 | 62  | 2014-12-31   | stories of our lives                        | stories of our lives                        | \$15,000          | \$0            | \$0             | 2014         | 1    | 1.0   |

1600 rows x 27 columns

Filter rows where 'movie' does not equal 'primary\_title' to manually explore the data

```
non_exact_matches = merged_df[merged_df['movie'] != merged_df['primary_title']]
```

```
# Set display options to show all rows
pd.set_option('display.max_rows', None)
```

```
non_exact_matches
```

10/26

|     |    |            |                                     |                                   |              |               |               |      |
|-----|----|------------|-------------------------------------|-----------------------------------|--------------|---------------|---------------|------|
| 447 | 41 | 2014-10-10 | the judge                           | the judgment                      | \$50,000,000 | \$47,119,388  | \$76,119,388  | 2014 |
| 464 | 4  | 2013-03-15 | upside down                         | upside down                       | \$50,000,000 | \$102,118     | \$26,387,039  | 2013 |
| 473 | 42 | 2016-12-23 | silence                             | silenced                          | \$46,500,000 | \$7,100,177   | \$23,726,626  | 2016 |
| 495 | 35 | 2013-06-28 | the heat                            | the east                          | \$43,000,000 | \$159,581,587 | \$229,727,774 | 2013 |
| 496 | 40 | 2018-11-09 | the girl in the spider's web        | the girl in the spider's web      | \$43,000,000 | \$14,828,555  | \$34,983,342  | 2018 |
| 511 | 81 | 2019-05-17 | john wick: chapter 3 - parable      | john wick: chapter 3 - parable    | \$40,000,000 | \$141,744,320 | \$256,498,033 | 2019 |
| 515 | 94 | 2017-02-10 | john wick: chapter two              | john wick: chapter 2              | \$40,000,000 | \$92,029,184  | \$171,350,009 | 2017 |
| 544 | 70 | 2011-12-25 | extremely loud and incredibly close | extremely loud & incredibly close | \$40,000,000 | \$31,847,881  | \$55,247,881  | 2011 |
| 560 | 30 | 2016-11-11 | billy lynn's long halftime walk     | billy lynn's long halftime walk   | \$40,000,000 | \$1,738,477   | \$30,230,402  | 2016 |
| 562 | 41 | 2014-12-31 | dragon nest warriors' dawn          | dragon nest: warriors' dawn       | \$40,000,000 | \$0           | \$734,423     | 2014 |
| 563 | 42 | 2018-12-31 | the crow                            | the row                           | \$40,000,000 | \$0           | \$0           | 2018 |
| 588 | 6  | 2011-02-11 | gnomeo and juliet                   | gnomeo & juliet                   | \$36,000,000 | \$99,967,670  | \$193,737,977 | 2011 |
| 619 | 3  | 2016-09-16 | bridget jones's baby                | bridget jones's baby              | \$35,000,000 | \$24,139,805  | \$205,822,688 | 2016 |
| 668 | 65 | 2017-08-18 | the hitman's bodyguard              | the hitman's bodyguard            | \$30,000,000 | \$75,468,583  | \$172,778,667 | 2017 |
| 708 | 79 | 2011-04-29 | hoodwinked too: hood vs. evil       | hoodwinked too! hood vs. evil     | \$30,000,000 | \$10,143,779  | \$23,353,111  | 2011 |
| 710 | 84 | 2017-02-03 | the space between us                | the space between                 | \$30,000,000 | \$7,885,294   | \$16,481,405  | 2017 |
| 742 | 19 | 2016-11-23 | rules don't apply                   | rules don't apply                 | \$26,700,000 | \$3,652,206   | \$3,871,448   | 2016 |
| 758 | 86 | 2017-01-27 | a dog's purpose                     | a dog's purpose                   | \$25,000,000 | \$64,321,890  | \$203,671,625 | 2017 |
| 853 | 39 | 2017-11-17 | the star                            | the stray                         | \$20,000,000 | \$40,847,995  | \$62,758,010  | 2017 |
| 861 | 61 | 2013-04-12 | scary movie v                       | scary movie 5                     | \$20,000,000 | \$32,015,787  | \$78,613,981  | 2013 |
| 865 | 80 | 2010-09-17 | alpha and omega 3d                  | alpha and omega                   | \$20,000,000 | \$25,107,267  | \$48,958,353  | 2010 |
| 871 | 5  | 2012-10-26 | silent hill: revelation 3d          | silent hill: revelation           | \$20,000,000 | \$17,530,219  | \$55,975,672  | 2012 |
| 872 | 7  | 2017-03-31 | the zookeeper's wife                | the zookeeper's wife              | \$20,000,000 | \$17,445,186  | \$26,308,749  | 2017 |
| 902 | 11 | 2015-05-29 | survivor                            | survivors                         | \$20,000,000 | \$0           | \$1,703,281   | 2015 |
| 921 | 64 | 2014-05-09 | neighbors                           | neighbours                        | \$18,000,000 | \$150,086,800 | \$270,944,428 | 2014 |
| 940 | 40 | 2017-03-17 | t2: trainspotting                   | t2 trainspotting                  | \$18,000,000 | \$2,402,004   | \$42,091,262  | 2017 |
| 946 | 53 | 2014-08-22 | the prince                          | the principle                     | \$18,000,000 | \$0           | \$0           | 2014 |
| 974 | 97 | 2010-11-26 | the king's speech                   | the king's speech                 | \$15,000,000 | \$138,797,449 | \$430,821,168 | 2010 |
| 988 | 39 | 2018-11-23 | the favourite                       | the favorite                      | \$15,000,000 | \$34,366,783  | \$94,113,929  | 2018 |

|      |    |            |                                                 |                                                   |              |              |               |      |
|------|----|------------|-------------------------------------------------|---------------------------------------------------|--------------|--------------|---------------|------|
| 1001 | 13 | 2013-01-04 | promised land                                   | promise land                                      | \$15,000,000 | \$7,597,898  | \$12,394,562  | 2013 |
| 1065 | 77 | 2016-08-12 | hell or high water                              | hell or high waters                               | \$12,000,000 | \$27,007,844 | \$37,584,304  | 2016 |
| 1085 | 55 | 2015-10-30 | dancin' it's on                                 | dancin': it's on!                                 | \$12,000,000 | \$0          | \$0           | 2015 |
| 1105 | 51 | 2015-06-05 | insidious chapter 3                             | insidious: chapter 3                              | \$10,000,000 | \$52,218,558 | \$120,453,155 | 2015 |
| 1144 | 69 | 2016-03-30 | everybody wants some                            | everybody wants some!!                            | \$10,000,000 | \$3,400,278  | \$5,437,126   | 2016 |
| 1160 | 33 | 2010-12-31 | the reef                                        | the tree                                          | \$10,000,000 | \$0          | \$15,037,867  | 2010 |
| 1240 | 99 | 2013-05-31 | the east                                        | the heat                                          | \$6,500,000  | \$2,274,649  | \$3,027,956   | 2013 |
| 1242 | 99 | 2013-05-31 | the east                                        | the past                                          | \$6,500,000  | \$2,274,649  | \$3,027,956   | 2013 |
| 1266 | 53 | 2013-09-13 | insidious chapter 2                             | insidious: chapter 2                              | \$5,000,000  | \$83,586,447 | \$161,921,515 | 2013 |
| 1270 | 56 | 2015-09-11 | the visit                                       | the visitor                                       | \$5,000,000  | \$65,206,105 | \$98,677,816  | 2015 |
| 1273 | 59 | 2012-10-19 | paranormal activity 4                           | paranormal captivity                              | \$5,000,000  | \$53,900,335 | \$142,817,992 | 2012 |
| 1291 | 8  | 2016-07-15 | hillary's america: the secret history of the... | hillary's america: the secret history of the d... | \$5,000,000  | \$13,099,931 | \$13,099,931  | 2016 |
| 1294 | 20 | 2014-05-09 | moms' night out                                 | moms' night out                                   | \$5,000,000  | \$10,429,707 | \$10,537,341  | 2014 |
| 1319 | 18 | 2012-09-14 | barfi                                           | barfi!                                            | \$4,600,000  | \$2,804,874  | \$36,751,984  | 2012 |
| 1325 | 56 | 2014-11-07 | fugly                                           | fugly!                                            | \$4,500,000  | \$0          | \$0           | 2014 |
| 1349 | 94 | 2016-02-19 | the witch                                       | the witching                                      | \$3,500,000  | \$25,138,705 | \$40,454,520  | 2016 |
| 1385 | 72 | 2015-08-14 | amnesiac                                        | amnesia                                           | \$3,000,000  | \$0          | \$0           | 2015 |
| 1392 | 13 | 2012-07-13 | 2016: obama's america                           | 2016: obama's america                             | \$2,500,000  | \$33,349,941 | \$33,349,941  | 2012 |
| 1401 | 44 | 2013-06-21 | alien uprising                                  | alien rising                                      | \$2,500,000  | \$0          | \$0           | 2013 |
| 1432 | 79 | 2013-04-01 | stitches                                        | stitch                                            | \$2,000,000  | \$0          | \$63,555      | 2013 |
| 1476 | 77 | 2018-02-06 | blood feast                                     | blood fest                                        | \$1,200,000  | \$8,708      | \$8,708       | 2018 |
| 1479 | 87 | 2015-02-03 | bleeding hearts                                 | bleeding heart                                    | \$1,200,000  | \$0          | \$0           | 2015 |
| 1516 | 48 | 2012-08-31 | for a good time, call                           | for a good time, call...                          | \$850,000    | \$1,251,749  | \$1,386,088   | 2012 |
| 1518 | 54 | 2012-08-03 | celeste and jesse forever                       | celeste & jesse forever                           | \$840,000    | \$3,103,407  | \$3,787,689   | 2012 |
| 1565 | 33 | 2014-03-14 | the word                                        | the m word                                        | \$200,000    | \$3,648      | \$3,648       | 2014 |
| 1568 | 43 | 2011-09-23 | weekend                                         | weekender                                         | \$190,000    | \$484,592    | \$1,577,585   | 2011 |
| 1572 | 60 | 2015-04-17 | antarctic edge: 70° south                       | antarctic edge: 70° south                         | \$150,000    | \$7,193      | \$7,193       | 2015 |
| 1581 | 93 | 2014-12-31 | dude, where's my dog                            | dude, where's my dog?!                            | \$100,000    | \$0          | \$0           | 2014 |
| 1585 | 6  | 2011-12-31 | absentia                                        | absent                                            | \$70,000     | \$0          | \$8,555       | 2011 |

84 rows × 27 columns

```
# Reset display options to default
pd.reset_option('display.max_rows')
```

Let's remove movies with poorly matched titles by index

```
poor_match = [97, 145, 447, 495, 563, 853, 946, 1160, 1240, 1242]
merged_df = merged_df.drop(poor_match)
```

```
merged_df.columns
```

```
Index(['id', 'release_date', 'movie', 'primary_title', 'production_budget',
       'domestic_gross', 'worldwide_gross', 'release_year', 'year', 'title',
       'imdb_movie_id', 'average_rating', 'numvotes', 'start_year',
       'runtime_minutes', 'genres', 'director1', 'director2', 'actress1',
       'actress2', 'actress3', 'actor1', 'actor2', 'actor3', 'actor4',
       'producer1', 'producer2'],
      dtype='object')
```

We can remove redundant columns

```
col_to_remove = ['id', 'movie', 'release_year', 'year', 'title', 'start_year', 'runtime_minutes']
merged_df = merged_df.drop(col_to_remove, axis=1)
```

```
merged_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 1590 entries, 0 to 1599
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   release_date          1590 non-null  datetime64[ns]
1   primary_title          1590 non-null  object
2   production_budget      1590 non-null  object
3   domestic_gross         1590 non-null  object
4   worldwide_gross        1590 non-null  object
5   imdb_movie_id         1590 non-null  object
6   average_rating         1590 non-null  float64
7   numvotes              1590 non-null  int64
8   genres                 1590 non-null  object
9   director1              1589 non-null  object
10  director2              139 non-null   object
11  actress1               1383 non-null  object
12  actress2               730 non-null   object
13  actress3               177 non-null   object
14  actor1                 1534 non-null  object
15  actor2                 1379 non-null  object
16  actor3                 825 non-null   object
17  actor4                 179 non-null   object
18  producer1              1380 non-null  object
19  producer2              997 non-null   object
dtypes: datetime64[ns](1), float64(1), int64(1), object(17)
memory usage: 260.9+ KB
```

```
# convert string columns to numeric
for col in ['production_budget', 'domestic_gross', 'worldwide_gross']:
    merged_df[col] = merged_df[col].str.replace('$', '', regex=False)
    merged_df[col] = merged_df[col].str.replace(',', '', regex=False)
    merged_df[col] = pd.to_numeric(merged_df[col])
```

```
# Calculate profit
merged_df['profit'] = merged_df['worldwide_gross'] - merged_df['production_budget']
```

```
# Calculate ROI
merged_df['roi'] = ((merged_df['profit'] - merged_df['production_budget']) / merged_df['production_budget']) * 100
```

```
merged_df.head()
```




|   | release_date | primary_title                               | production_budget | domestic_gross | worldwide_gross | imdb_movie_id | average_rating | numvotes |   |
|---|--------------|---------------------------------------------|-------------------|----------------|-----------------|---------------|----------------|----------|---|
| 0 | 2011-05-20   | pirates of the caribbean: on stranger tides | 410600000         | 241063875      | 1045663875      | tt1298650     | 6.6            | 447624   | A |
| 1 | 2019-06-07   | dark phoenix                                | 350000000         | 42762350       | 149762350       | tt6565702     | 6.0            | 24451    |   |
| 2 | 2015-05-01   | avengers: age of ultron                     | 330600000         | 459005868      | 1403013963      | tt2395427     | 7.3            | 665594   |   |
| 3 | 2018-04-27   | avengers: infinity war                      | 300000000         | 678815482      | 2048134200      | tt4154756     | 8.5            | 670926   |   |
| 4 | 2017-11-17   | justice league                              | 300000000         | 229024295      | 655945209       | tt0974015     | 6.5            | 329135   | A |

5 rows x 22 columns

Split the genres column into multiple columns

```
merged_df[['genre1', 'genre2', 'genre3']] = merged_df['genres'].str.split(',', expand=True)# expand=True: Ensures that the resul
```

```
merged_df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
Index: 1590 entries, 0 to 1599
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype
---  -
0   release_date          1590 non-null   datetime64[ns]
1   primary_title         1590 non-null   object
2   production_budget     1590 non-null   int64
3   domestic_gross        1590 non-null   int64
4   worldwide_gross       1590 non-null   int64
5   imdb_movie_id         1590 non-null   object
6   average_rating        1590 non-null   float64
7   numvotes              1590 non-null   int64
8   genres                1590 non-null   object
9   director1            1589 non-null   object
10  director2             139 non-null    object
11  actress1              1383 non-null   object
12  actress2              730 non-null    object
13  actress3              177 non-null    object
14  actor1                1534 non-null   object
15  actor2                1379 non-null   object
16  actor3                825 non-null    object
17  actor4                179 non-null    object
18  producer1            1380 non-null   object
19  producer2            997 non-null    object
20  profit               1590 non-null   int64
21  roi                  1590 non-null   float64
22  genre1               1590 non-null   object
23  genre2               1408 non-null   object
24  genre3               1046 non-null   object
dtypes: datetime64[ns](1), float64(2), int64(5), object(17)
memory usage: 323.0+ KB
```

## ✓ Genre Analysis

We can set the float format to display up to a specific number of decimal places

```
pd.set_option('display.float_format', lambda x: '%.2f' % x if x != 0 else 0)
```

Let's melting the genre columns and group by genre and calculate metrics

```
# Melting the genre columns
melted_genres_df = merged_df.melt(id_vars=['profit', 'roi', 'average_rating', 'numvotes', 'production_budget'],
                                value_vars=['genre1', 'genre2', 'genre3'],
                                value_name='genre').dropna(subset=['genre'])
```

```
# Group by genre and calculate metrics
genre_analysis = melted_genres_df.groupby('genre').agg(
    avg_budget=('production_budget', 'mean'),
    avg_profit=('profit', 'mean'),
    avg_roi=('roi', 'mean'),
    avg_rating=('average_rating', 'mean'),
    total_votes=('numvotes', 'sum')
).reset_index()

genre_analysis = genre_analysis.sort_values(by=['avg_roi', 'avg_rating'], ascending=[False, False])

genre_analysis
```



|    | genre       | avg_budget   | avg_profit   | avg_roi | avg_rating | total_votes |
|----|-------------|--------------|--------------|---------|------------|-------------|
| 14 | Mystery     | 25125550.38  | 71788590.80  | 782.92  | 6.10       | 16594459    |
| 11 | Horror      | 18517602.59  | 53785183.95  | 701.17  | 5.45       | 12523815    |
| 18 | Thriller    | 33116149.45  | 85528899.22  | 432.83  | 6.00       | 32491608    |
| 12 | Music       | 17200754.72  | 61393441.79  | 182.88  | 6.33       | 3572520     |
| 3  | Biography   | 25500143.88  | 56630786.06  | 158.11  | 7.01       | 15425091    |
| 2  | Animation   | 93888392.86  | 262007659.49 | 157.61  | 6.49       | 13066701    |
| 16 | Sci-Fi      | 92910000.00  | 243146307.59 | 156.71  | 6.41       | 37240239    |
| 15 | Romance     | 22858128.08  | 46024403.09  | 148.43  | 6.30       | 16168782    |
| 7  | Drama       | 27205330.64  | 49728696.69  | 119.25  | 6.53       | 75170946    |
| 9  | Fantasy     | 85167537.31  | 168263427.75 | 118.22  | 6.11       | 21040550    |
| 1  | Adventure   | 105360923.48 | 245259640.20 | 109.90  | 6.44       | 74095851    |
| 10 | History     | 33772727.27  | 58648039.41  | 108.04  | 6.86       | 4561316     |
| 4  | Comedy      | 40577748.64  | 92092923.88  | 106.71  | 6.19       | 50991232    |
| 6  | Documentary | 5360756.76   | 11935614.24  | 82.96   | 6.62       | 341119      |
| 8  | Family      | 65006000.00  | 120873862.53 | 67.79   | 6.05       | 7224888     |
| 0  | Action      | 79215253.16  | 160527513.80 | 62.33   | 6.25       | 83036951    |
| 5  | Crime       | 33050000.00  | 51401423.40  | 17.89   | 6.31       | 27679957    |
| 13 | Musical     | 40035000.00  | 144864392.90 | 12.86   | 5.48       | 615712      |
| 17 | Sport       | 24084558.82  | 27435814.06  | 9.99    | 6.89       | 2395967     |
| 20 | Western     | 55983333.33  | 47182244.17  | -41.26  | 6.30       | 2158691     |
| 19 | War         | 27416666.67  | 27183357.72  | -66.51  | 6.45       | 1290577     |

Now let's group by genre combinations and sort by the highest ROI

```
# Grouping by genre combinations and calculating metrics
genre_combinations_analysis = merged_df.groupby('genres').agg(
    count=('genres', 'size'), # Count the number of occurrences of each genre combination
    avg_budget=('production_budget', 'mean'),
    avg_profit=('profit', 'mean'),
    avg_roi=('roi', 'mean'),
    avg_rating=('average_rating', 'mean'),
    total_votes=('numvotes', 'sum')
).reset_index()

# Filter for genre combinations with at least 30 occurrences
genre_combinations_analysis_filtered = genre_combinations_analysis[genre_combinations_analysis['count'] >= 30]

# BY ROI
# Sort by average ROI to get the top ROI combinations
genre_combinations_analysis_filtered = genre_combinations_analysis_filtered.sort_values(by='avg_roi', ascending=False)

# Display the top 20 genre combinations based on average ROI
genre_combinations_analysis_filtered.head(20)
```





|     | genres                     | count | avg_budget   | avg_profit   | avg_roi | avg_rating | total_votes |
|-----|----------------------------|-------|--------------|--------------|---------|------------|-------------|
| 225 | Horror,Mystery,Thriller    | 33    | 9721969.70   | 97892046.21  | 2614.19 | 5.53       | 2662668     |
| 203 | Drama,Romance              | 35    | 22846857.14  | 54608913.46  | 270.25  | 6.82       | 3008082     |
| 63  | Adventure,Animation,Comedy | 72    | 100090277.78 | 292000004.51 | 192.66  | 6.41       | 8178889     |
| 12  | Action,Adventure,Sci-Fi    | 55    | 169716363.64 | 467118871.35 | 166.97  | 6.73       | 22181507    |
| 132 | Comedy,Drama,Romance       | 58    | 18023965.52  | 29467109.05  | 122.75  | 6.38       | 4758146     |
| 152 | Comedy,Romance             | 44    | 25565909.09  | 56749376.34  | 119.30  | 5.85       | 3215413     |
| 172 | Drama                      | 61    | 11451672.13  | 14978908.18  | 113.12  | 6.70       | 2879926     |
| 124 | Comedy,Drama               | 52    | 16168461.54  | 27646498.40  | 94.45   | 6.44       | 2985850     |
| 8   | Action,Adventure,Fantasy   | 33    | 149109090.91 | 247778337.39 | 48.95   | 6.12       | 7794508     |
| 21  | Action,Comedy,Crime        | 33    | 39930000.00  | 58148700.58  | 44.70   | 6.03       | 3794183     |
| 117 | Comedy                     | 63    | 29061190.48  | 41244930.44  | 24.59   | 5.54       | 3906958     |
| 32  | Action,Crime,Drama         | 39    | 27548717.95  | 25689928.59  | 2.77    | 6.35       | 3922818     |
| 6   | Action,Adventure,Drama     | 30    | 88153333.33  | 84898672.47  | -17.70  | 5.96       | 3590583     |

```
# Grouping by genre combinations and calculating metrics
genre_combinations_analysis = merged_df.groupby('genres').agg(
    count=('genres', 'size'), # Count the number of occurrences of each genre combination
    avg_budget=('production_budget', 'mean'),
    avg_profit=('profit', 'mean'),
    avg_roi=('roi', 'mean'),
    avg_rating=('average_rating', 'mean'),
    total_votes=('numvotes', 'sum')
).reset_index()

# Filter for genre combinations with at least 30 occurrences
genre_combinations_analysis_filtered = genre_combinations_analysis[genre_combinations_analysis['count'] >= 30]

# BY ROI
# Sort by average ROI to get the top ROI combinations
genre_combinations_analysis_filtered = genre_combinations_analysis_filtered.sort_values(by='avg_roi', ascending=False)

# Display the top 20 genre combinations based on average ROI
genre_combinations_analysis_filtered.head(20)
```



|     | genres                     | count | avg_budget   | avg_profit   | avg_roi | avg_rating | total_votes |
|-----|----------------------------|-------|--------------|--------------|---------|------------|-------------|
| 225 | Horror,Mystery,Thriller    | 33    | 9721969.70   | 97892046.21  | 2614.19 | 5.53       | 2662668     |
| 203 | Drama,Romance              | 35    | 22846857.14  | 54608913.46  | 270.25  | 6.82       | 3008082     |
| 63  | Adventure,Animation,Comedy | 72    | 100090277.78 | 292000004.51 | 192.66  | 6.41       | 8178889     |
| 12  | Action,Adventure,Sci-Fi    | 55    | 169716363.64 | 467118871.35 | 166.97  | 6.73       | 22181507    |
| 132 | Comedy,Drama,Romance       | 58    | 18023965.52  | 29467109.05  | 122.75  | 6.38       | 4758146     |
| 152 | Comedy,Romance             | 44    | 25565909.09  | 56749376.34  | 119.30  | 5.85       | 3215413     |
| 172 | Drama                      | 61    | 11451672.13  | 14978908.18  | 113.12  | 6.70       | 2879926     |
| 124 | Comedy,Drama               | 52    | 16168461.54  | 27646498.40  | 94.45   | 6.44       | 2985850     |
| 8   | Action,Adventure,Fantasy   | 33    | 149109090.91 | 247778337.39 | 48.95   | 6.12       | 7794508     |
| 21  | Action,Comedy,Crime        | 33    | 39930000.00  | 58148700.58  | 44.70   | 6.03       | 3794183     |
| 117 | Comedy                     | 63    | 29061190.48  | 41244930.44  | 24.59   | 5.54       | 3906958     |
| 32  | Action,Crime,Drama         | 39    | 27548717.95  | 25689928.59  | 2.77    | 6.35       | 3922818     |
| 6   | Action,Adventure,Drama     | 30    | 88153333.33  | 84898672.47  | -17.70  | 5.96       | 3590583     |

Let's include some visualizations

```
import matplotlib.pyplot as plt
import seaborn as sns

# Set modern and professional style
```

```

sns.set_style("whitegrid")
plt.figure(figsize=(12, 8))

# Sort by 'avg_roi' to get the top 10 genre combinations
top_10_genres_by_roi = genre_combinations_analysis_filtered.nlargest(10, 'avg_roi')

# Create the bar plot with dodgerblue color
barplot = sns.barplot(x='avg_roi', y='genres', data=top_10_genres_by_roi, color='dodgerblue', edgecolor='black')

# Customize title and labels for a professional look
barplot.set_title('Top 10 Genre Combinations by ROI', fontsize=18, weight='bold', color='navy', pad=20)
barplot.set_xlabel('Average ROI', fontsize=14, weight='bold')
barplot.set_ylabel('Genre Combinations', fontsize=14, weight='bold')

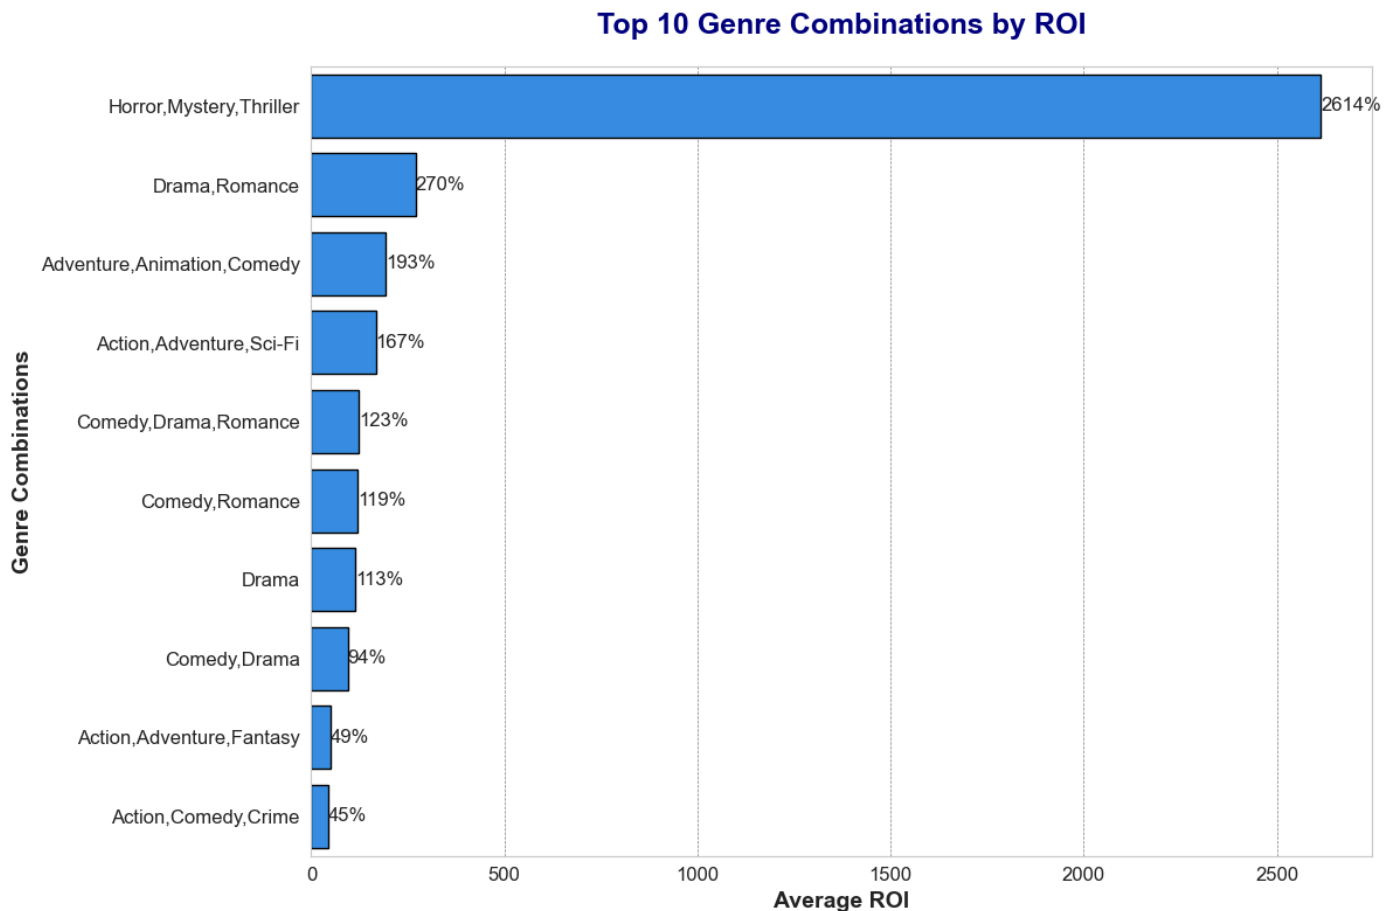
# Add grid for better readability
barplot.xaxis.grid(True, color='gray', linestyle='--', linewidth=0.5)

# Customize ticks
barplot.tick_params(labelsize=12)

# Add value labels for each bar, rounded to integer
for i in barplot.containers:
    barplot.bar_label(i, fmt='%.0f%%', fontsize=12) # Rounded to integer

# Show the plot
plt.tight_layout()
plt.show()

```



Now we can sort by average profit

```
# Sort by average profit to get the top profitable combinations
genre_combinations_analysis_filtered = genre_combinations_analysis_filtered.sort_values(by='avg_profit', ascending=False)

# Display the top 20 genre combinations based on average profit
genre_combinations_analysis_filtered.head(20)
```



|     | genres                     | count | avg_budget   | avg_profit   | avg_roi | avg_rating | total_votes |
|-----|----------------------------|-------|--------------|--------------|---------|------------|-------------|
| 12  | Action,Adventure,Sci-Fi    | 55    | 169716363.64 | 467118871.35 | 166.97  | 6.73       | 22181507    |
| 63  | Adventure,Animation,Comedy | 72    | 100090277.78 | 292000004.51 | 192.66  | 6.41       | 8178889     |
| 8   | Action,Adventure,Fantasy   | 33    | 149109090.91 | 247778337.39 | 48.95   | 6.12       | 7794508     |
| 225 | Horror,Mystery,Thriller    | 33    | 9721969.70   | 97892046.21  | 2614.19 | 5.53       | 2662668     |
| 6   | Action,Adventure,Drama     | 30    | 88153333.33  | 84898672.47  | -17.70  | 5.96       | 3590583     |
| 21  | Action,Comedy,Crime        | 33    | 39930000.00  | 58148700.58  | 44.70   | 6.03       | 3794183     |
| 152 | Comedy,Romance             | 44    | 25565909.09  | 56749376.34  | 119.30  | 5.85       | 3215413     |
| 203 | Drama,Romance              | 35    | 22846857.14  | 54608913.46  | 270.25  | 6.82       | 3008082     |
| 117 | Comedy                     | 63    | 29061190.48  | 41244930.44  | 24.59   | 5.54       | 3906958     |
| 132 | Comedy,Drama,Romance       | 58    | 18023965.52  | 29467109.05  | 122.75  | 6.38       | 4758146     |
| 124 | Comedy,Drama               | 52    | 16168461.54  | 27646498.40  | 94.45   | 6.44       | 2985850     |
| 32  | Action,Crime,Drama         | 39    | 27548717.95  | 25689928.59  | 2.77    | 6.35       | 3922818     |
| 172 | Drama                      | 61    | 11451672.13  | 14978908.18  | 113.12  | 6.70       | 2879926     |

```
import matplotlib.pyplot as plt
import seaborn as sns
import matplotlib.ticker as ticker

# Function to format numbers as currency with millions
def format_currency_millions(value):
    return f"${value / 1_000_000:.1f}M" # Format as millions with one decimal place

# Set modern and professional style
sns.set_style("whitegrid")
plt.figure(figsize=(12, 8))

# Sort by 'avg_profit' to get the top 10 genre combinations
top_10_genres_by_profit = genre_combinations_analysis_filtered.nlargest(10, 'avg_profit')

# Create the bar plot with dodgerblue color
barplot = sns.barplot(x='avg_profit', y='genres', data=top_10_genres_by_profit, color='dodgerblue', edgecolor='black')

# Customize title and labels for a professional look
barplot.set_title('Top 10 Genre Combinations by Average Profit', fontsize=18, weight='bold', color='navy', pad=20)
barplot.set_xlabel('Average Profit', fontsize=14, weight='bold') # Proper x-axis title
barplot.set_ylabel('Genre Combinations', fontsize=14, weight='bold')

# Add grid for better readability
barplot.xaxis.grid(True, color='gray', linestyle='--', linewidth=0.5)

# Remove x-axis tick values but keep the grid
barplot.xaxis.set_major_locator(ticker.MaxNLocator(integer=True)) # Ensures grid lines are integer values
barplot.set_xticklabels([]) # Remove x-tick labels


# Customize tick parameters
barplot.tick_params(labelsize=12)

# Add value labels for each bar formatted as currency in millions, with adjusted positioning
for p in barplot.patches:
    barplot.annotate(format_currency_millions(p.get_x() + p.get_width()), # Format the label
                    (p.get_x() + p.get_width(), p.get_y() + p.get_height() / 2), # Position the label
                    ha='left', va='center', fontsize=12, color='black',
                    xytext=(5, 0), textcoords='offset points') # Offset from the bar

# Show the plot
plt.tight_layout()
plt.show()
```



## Top 10 Genre Combinations by Average Profit



### ✓ Horror, Mystery, Thriller

### ✓ Find the most successful producers

Find best producers for Horror, Mystery, or Thriller by the weighted score of (avg\_roi \* movie\_count). We can include all movies tagged as either Horror, Mystery, or Thriller:

```
# Create a mask for the specified genres
mask_genres_hmt = merged_df['genres'] == 'Horror,Mystery,Thriller'
mask_separate_genre_hmt = merged_df['genre1'].isin(['Horror', 'Mystery', 'Thriller']) | \
    merged_df['genre2'].isin(['Horror', 'Mystery', 'Thriller']) | \
    merged_df['genre3'].isin(['Horror', 'Mystery', 'Thriller'])
```

```
# Combine the masks
hmt_all_df = merged_df[mask_genres_hmt | mask_separate_genre_hmt]
```

If we want to include only specific genre category 'Horror,Mystery,Thriller'

```
# Filter the DataFrame for the specified genres
hmt_df = merged_df[merged_df['genres'] == 'Horror,Mystery,Thriller']
```

```
# Include separate genres
# Melt the producers into a single column
producers = pd.melt(hmt_all_df, id_vars=['profit', 'roi'],
    value_vars=['producer1', 'producer2'],
    var_name='producer_type', value_name='producer_name')
```

```
# Remove rows with NaN producers
```

```

producers = producers.dropna(subset=['producer_name'])

# Group by producer name and calculate average profit and movie count
producer_analysis = producers.groupby('producer_name').agg(
    avg_profit=('profit', 'mean'),
    avg_roi=('roi', 'mean'),
    movie_count=('roi', 'count')
).reset_index()

# Calculate a weighted score
producer_analysis['total_profit'] = producer_analysis['avg_profit'] * producer_analysis['movie_count']

# Sort by score to find the best producers
best_producers = producer_analysis.sort_values(by='total_profit', ascending=False)

# Display the top producers
best_producers.head(10)

```



|     | producer_name    | avg_profit   | avg_roi | movie_count | total_profit  |
|-----|------------------|--------------|---------|-------------|---------------|
| 368 | Neal H. Moritz   | 862008239.75 | 372.68  | 4           | 3448032959.00 |
| 211 | Jason Blum       | 85844562.28  | 2538.74 | 36          | 3090404242.00 |
| 340 | Michael Fottrell | 939577505.00 | 399.13  | 3           | 2818732515.00 |
| 125 | David Heyman     | 709565120.50 | 499.49  | 2           | 1419130241.00 |
| 109 | Dana Brunetti    | 255373029.25 | 457.40  | 4           | 1021492117.00 |
| 334 | Michael Bay      | 112803038.22 | 1141.24 | 9           | 1015227344.00 |
| 338 | Michael De Luca  | 202912027.20 | 344.20  | 5           | 1014560136.00 |
| 208 | James Wan        | 190578607.60 | 2169.43 | 5           | 952893038.00  |
| 44  | Barbara Broccoli | 910526981.00 | 355.26  | 1           | 910526981.00  |
| 119 | David Barron     | 453404490.00 | 293.65  | 2           | 906808980.00  |

```

# Include only specific genre category 'Horror,Mystery,Thriller'

# Melt the producers into a single column
producers = pd.melt(hmt_df, id_vars=['profit', 'roi'],
    value_vars=['producer1', 'producer2'],
    var_name='producer_type', value_name='producer_name')

# Remove rows with NaN producers
producers = producers.dropna(subset=['producer_name'])

# Group by producer name and calculate average profit and movie count
producer_analysis = producers.groupby('producer_name').agg(
    avg_profit=('profit', 'mean'),
    avg_roi=('roi', 'mean'),
    movie_count=('roi', 'count')
).reset_index()

# Calculate a weighted score
producer_analysis['total_profit'] = producer_analysis['avg_profit'] * producer_analysis['movie_count']

# Sort by score to find the best producers
best_producers = producer_analysis.sort_values(by='total_profit', ascending=False)

# Display the top producers
best_producers.head(10)

```



|    | producer_name     | avg_profit   | avg_roi | movie_count | total_profit  |
|----|-------------------|--------------|---------|-------------|---------------|
| 8  | Jason Blum        | 129855601.64 | 5012.65 | 14          | 1817978423.00 |
| 7  | James Wan         | 202271632.00 | 2017.76 | 4           | 809086528.00  |
| 19 | Peter Safran      | 270373892.50 | 2793.82 | 2           | 540747785.00  |
| 26 | Sean McKittrick   | 242289130.50 | 2989.21 | 2           | 484578261.00  |
| 30 | Tony DeRosa-Grund | 298000141.00 | 1390.00 | 1           | 298000141.00  |
| 20 | Rob Cowan         | 298000141.00 | 1390.00 | 1           | 298000141.00  |
| 28 | Steven Schneider  | 202039844.00 | 3940.80 | 1           | 202039844.00  |
| 18 | Oren Peli         | 156921515.00 | 3038.43 | 1           | 156921515.00  |
| 16 | Michael Bay       | 98300632.00  | 1866.01 | 1           | 98300632.00   |
| 14 | Marc Bienstock    | 93677816.00  | 1773.56 | 1           | 93677816.00   |

## Find the most successful directors

Find best directors for Horror, Mystery, or Thriller by the weighted score of (avg\_roi \* movie\_count).

Include separate genres

```
# Melt the directors into a single column
directors = pd.melt(hmt_all_df, id_vars=['profit', 'roi'],
                    value_vars=['director1', 'director2'],
                    var_name='director_type', value_name='director_name')

# Remove rows with NaN producers
directors = directors.dropna(subset=['director_name'])

# Group by director name and calculate average roi and movie count
director_analysis = directors.groupby('director_name').agg(
    avg_profit=('profit', 'mean'),
    avg_roi=('roi', 'mean'),
    movie_count=('roi', 'count')
).reset_index()

# Calculate a weighted score
director_analysis['total_profit'] = director_analysis['avg_profit'] * director_analysis['movie_count']

# Sort by score to find the best director
best_directors = director_analysis.sort_values(by='total_profit', ascending=False)

# Display the top directors
best_directors.head(10)
```



|     | director_name         | avg_profit   | avg_roi | movie_count | total_profit  |
|-----|-----------------------|--------------|---------|-------------|---------------|
| 150 | James Wan             | 594548150.00 | 1675.92 | 3           | 1783644450.00 |
| 310 | Sam Mendes            | 745073952.00 | 224.24  | 2           | 1490147904.00 |
| 188 | Justin Lin            | 567231949.00 | 298.72  | 2           | 1134463898.00 |
| 111 | F. Gary Gray          | 984846267.00 | 293.94  | 1           | 984846267.00  |
| 87  | David Yates           | 835431568.00 | 568.35  | 1           | 835431568.00  |
| 53  | Christopher Nolan     | 809439099.00 | 194.34  | 1           | 809439099.00  |
| 18  | Andy Muschietti       | 397776767.50 | 1290.02 | 2           | 795553535.00  |
| 52  | Christopher McQuarrie | 383413644.00 | 202.34  | 2           | 766827288.00  |
| 304 | Ruben Fleischer       | 737628605.00 | 535.89  | 1           | 737628605.00  |
| 253 | Neil Boulby           | 623008101.00 | 256.00  | 1           | 623008101.00  |

Include only specific genre category 'Horror,Mystery,Thriller'

```
# Melt the producers into a single column
directors = pd.melt(hmt_df, id_vars=['profit', 'roi'],
```

```

value_vars=['director1', 'director2'],
var_name='director_type', value_name='director_name')

# Remove rows with NaN producers
directors = directors.dropna(subset=['director_name'])

# Group by director name and calculate average roi and movie count
director_analysis = directors.groupby('director_name').agg(
    avg_profit=('profit', 'mean'),
    avg_roi=('roi', 'mean'),
    movie_count=('roi', 'count')
).reset_index()

# Calculate a weighted score
director_analysis['total_profit'] = director_analysis['avg_profit'] * director_analysis['movie_count']

# Sort by score to find the best director
best_directors = director_analysis.sort_values(by='total_profit', ascending=False)

# Display the top directors
best_directors.head(10)

```



|    | director_name      | avg_profit   | avg_roi | movie_count | total_profit |
|----|--------------------|--------------|---------|-------------|--------------|
| 18 | Jordan Peele       | 242289130.50 | 2989.21 | 2           | 484578261.00 |
| 14 | James Wan          | 227460828.00 | 2214.22 | 2           | 454921656.00 |
| 2  | Ariel Schulman     | 169928918.00 | 3298.58 | 2           | 339857836.00 |
| 13 | Henry Joost        | 169928918.00 | 3298.58 | 2           | 339857836.00 |
| 9  | David F. Sandberg  | 290384865.00 | 1835.90 | 1           | 290384865.00 |
| 17 | John R. Leonetti   | 250362920.00 | 3751.74 | 1           | 250362920.00 |
| 6  | Christopher Landon | 102957557.00 | 1959.15 | 2           | 205915114.00 |
| 0  | Adam Robitel       | 157885588.00 | 1478.86 | 1           | 157885588.00 |
| 25 | Scott Derrickson   | 71342212.50  | 1408.72 | 2           | 142684425.00 |
| 20 | Leigh Whannell     | 110453155.00 | 1004.53 | 1           | 110453155.00 |

## Find the most successful actresses

Find top 10 actresses based on score(total profit = avg\_profit \* movie\_count).

Include separate genres

```

# Combine all actress columns into a single column for analysis
actresses_df = hmt_all_df.melt(id_vars=['profit', 'roi'], value_vars=['actress1', 'actress2', 'actress3'],
    var_name='actress_rank', value_name='actress')

# Remove rows with missing actress values
actresses_df = actresses_df[actresses_df['actress'].notna()]

# Group by actress and calculate average profit, ROI, and total movies
best_actresses = actresses_df.groupby('actress').agg(
    avg_profit=('profit', 'mean'),
    avg_roi=('roi', 'mean'),
    movie_count=('actress', 'size') # Count how many movies each actress appeared in
).reset_index()

# Calculate a weighted score
best_actresses['total_profit'] = best_actresses['avg_profit'] * best_actresses['movie_count']

# Sort by avg_profit to find the best actresses
best_actresses = best_actresses.sort_values(by='total_profit', ascending=False)

best_actresses.head(10)

```



|     | actress              | avg_profit   | avg_roi | movie_count | total_profit  |
|-----|----------------------|--------------|---------|-------------|---------------|
| 341 | Michelle Rodriguez   | 371506279.60 | 182.62  | 5           | 1857531398.00 |
| 239 | Judi Dench           | 600724855.50 | 392.11  | 2           | 1201449711.00 |
| 363 | Naomie Harris        | 910526981.00 | 355.26  | 1           | 910526981.00  |
| 161 | Emma Watson          | 284635935.33 | 157.77  | 3           | 853907806.00  |
| 154 | Eloise Mumford       | 423674360.00 | 851.34  | 2           | 847348720.00  |
| 124 | Dakota Johnson       | 281517699.67 | 531.31  | 3           | 844553099.00  |
| 49  | Anne Hathaway        | 397903064.00 | 19.90   | 2           | 795806128.00  |
| 344 | Michelle Williams    | 737628605.00 | 535.89  | 1           | 737628605.00  |
| 453 | Sophia Lillis        | 662457969.00 | 1792.74 | 1           | 662457969.00  |
| 364 | Natalia Kaverznikova | 623008101.00 | 256.00  | 1           | 623008101.00  |

```
# Find top 10 actresses based on score(total profit = avg_profit * movie_count)
# Don't Include separate genres

# Combine all actress columns into a single column for analysis
actresses_df = hmt_df.melt(id_vars=['profit', 'roi'], value_vars=['actress1', 'actress2', 'actress3'],
                           var_name='actress_rank', value_name='actress')

# Remove rows with missing actress values
actresses_df = actresses_df[actresses_df['actress'].notna()]

# Group by actress and calculate average profit, ROI, and total movies
best_actresses = actresses_df.groupby('actress').agg(
    avg_profit=('profit', 'mean'),
    avg_roi=('roi', 'mean'),
    movie_count=('actress', 'size') # Count how many movies each actress appeared in
).reset_index()

# Calculate a weighted score
best_actresses['total_profit'] = best_actresses['avg_profit'] * best_actresses['movie_count']

# Sort by avg_profit to find the best actresses
best_actresses = best_actresses.sort_values(by='total_profit', ascending=False)

best_actresses.head(10)
```



|    | actress          | avg_profit   | avg_roi | movie_count | total_profit |
|----|------------------|--------------|---------|-------------|--------------|
| 35 | Lin Shaye        | 157403551.50 | 2258.64 | 2           | 314807103.00 |
| 55 | Vera Farmiga     | 298000141.00 | 1390.00 | 1           | 298000141.00 |
| 34 | Lili Taylor      | 298000141.00 | 1390.00 | 1           | 298000141.00 |
| 50 | Samara Lee       | 290384865.00 | 1835.90 | 1           | 290384865.00 |
| 38 | Miranda Otto     | 290384865.00 | 1835.90 | 1           | 290384865.00 |
| 2  | Allison Williams | 250367951.00 | 4907.36 | 1           | 250367951.00 |
| 14 | Catherine Keener | 250367951.00 | 4907.36 | 1           | 250367951.00 |
| 0  | Alfre Woodard    | 250362920.00 | 3751.74 | 1           | 250362920.00 |
| 5  | Annabelle Wallis | 250362920.00 | 3751.74 | 1           | 250362920.00 |
| 36 | Lupita Nyong'o   | 234210310.00 | 1071.05 | 1           | 234210310.00 |

## ✓ Find the most successful actors

Find top 10 actors based on score(total profit = avg\_profit \* movie\_count).

Include separate genres

```
# Combine all actor columns into a single column for analysis
actors_df = hmt_all_df.melt(id_vars=['profit', 'roi'], value_vars=['actor1', 'actor2', 'actor3', 'actor4'],
                             var_name='actor_rank', value_name='actor')
```



```
# Remove rows with missing actor values
actors_df = actors_df[actors_df['actor'].notna()]

# Group by actress and calculate average profit, ROI, and total movies
best_actors = actors_df.groupby('actor').agg(
    avg_profit=('profit', 'mean'),
    avg_roi=('roi', 'mean'),
    movie_count=('actor', 'size') # Count how many movies each actor appeared in
).reset_index()

# Calculate a weighted score
best_actors['total_profit'] = best_actors['avg_profit'] * best_actors['movie_count']

# Sort by avg_profit to find the best actor
best_actors = best_actors.sort_values(by='total_profit', ascending=False)

best_actors.head(10)
```



|     | actor            | avg_profit   | avg_roi | movie_count | total_profit  |
|-----|------------------|--------------|---------|-------------|---------------|
| 774 | Vin Diesel       | 741613263.60 | 339.33  | 5           | 3708066318.00 |
| 214 | Dwayne Johnson   | 611676037.83 | 286.68  | 6           | 3670056227.00 |
| 336 | Jason Statham    | 282868960.91 | 88.57   | 11          | 3111558570.00 |
| 592 | Paul Walker      | 821062230.67 | 398.92  | 3           | 2463186692.00 |
| 156 | Daniel Craig     | 407791010.00 | 99.43   | 4           | 1631164040.00 |
| 748 | Tom Hardy        | 535840171.67 | 306.03  | 3           | 1607520515.00 |
| 256 | Gary Oldman      | 207737957.20 | 123.45  | 5           | 1038689786.00 |
| 339 | Javier Bardem    | 323022463.67 | 127.02  | 3           | 969067391.00  |
| 161 | Daniel Radcliffe | 313503944.33 | 368.62  | 3           | 940511833.00  |
| 745 | Tom Cruise       | 288924592.33 | 157.08  | 3           | 866773777.00  |

```
# Find top 10 actors based on score(total profit = avg_profit * movie_count)
# Don't Include separate genres

# Combine all actor columns into a single column for analysis
actors_df = hmt_df.melt(id_vars=['profit', 'roi'], value_vars=['actor1', 'actor2', 'actor3', 'actor4'],
    var_name='actor_rank', value_name='actor')

# Remove rows with missing actor values
actors_df = actors_df[actors_df['actor'].notna()]

# Group by actress and calculate average profit, ROI, and total movies
best_actors = actors_df.groupby('actor').agg(
    avg_profit=('profit', 'mean'),
    avg_roi=('roi', 'mean'),
    movie_count=('actor', 'size') # Count how many movies each actor appeared in
).reset_index()

# Calculate a weighted score
best_actors['total_profit'] = best_actors['avg_profit'] * best_actors['movie_count']

# Sort by avg_profit to find the best actor
best_actors = best_actors.sort_values(by='total_profit', ascending=False)

best_actors.head(10)
```



|    | actor            | avg_profit   | avg_roi | movie_count | total_profit |
|----|------------------|--------------|---------|-------------|--------------|
| 50 | Patrick Wilson   | 227460828.00 | 2214.22 | 2           | 454921656.00 |
| 55 | Ron Livingston   | 298000141.00 | 1390.00 | 1           | 298000141.00 |
| 2  | Anthony LaPaglia | 290384865.00 | 1835.90 | 1           | 290384865.00 |
| 5  | Brad Greenquist  | 290384865.00 | 1835.90 | 1           | 290384865.00 |
| 1  | Angus Sampson    | 134169371.50 | 1241.69 | 2           | 268338743.00 |
| 43 | Leigh Whannell   | 134169371.50 | 1241.69 | 2           | 268338743.00 |
| 15 | Daniel Kaluuya   | 250367951.00 | 4907.36 | 1           | 250367951.00 |
| 6  | Bradley Whitford | 250367951.00 | 4907.36 | 1           | 250367951.00 |
| 65 | Ward Horton      | 250362920.00 | 3751.74 | 1           | 250362920.00 |
| 63 | Tony Amendola    | 250362920.00 | 3751.74 | 1           | 250362920.00 |

Start coding or [generate](#) with AI.

## ✓ Action, Adventure, Sci-Fi (maximize profit)

### ✓ Find the most successful producers

Find best producers for Action, Adventure, or Sci-Fi by the weighted score of (avg\_profit \* movie\_count).

Include either genre

```
aasf_all_df = merged_df[
    (merged_df['genre1'].isin(['Action', 'Adventure', 'Sci-Fi'])) |
    (merged_df['genre2'].isin(['Action', 'Adventure', 'Sci-Fi'])) |
    (merged_df['genre3'].isin(['Action', 'Adventure', 'Sci-Fi']))
]

# To include only specific genre category 'Action,Adventure,Sci-Fi'
# Filter the DataFrame for the specified genres
aasf_df = merged_df[merged_df['genres'] == 'Action,Adventure,Sci-Fi']

# Include separate genres

# Melt the producers into a single column
producers = pd.melt(aasf_all_df, id_vars=['profit'],
                    value_vars=['producer1', 'producer2'],
                    var_name='producer_type', value_name='producer_name')

# Remove rows with NaN producers
producers = producers.dropna(subset=['producer_name'])

# Group by producer name and calculate average profit and movie count
producer_analysis = producers.groupby('producer_name').agg(
    avg_profit=('profit', 'mean'),
    movie_count=('profit', 'count')
).reset_index()

# Calculate a weighted score
producer_analysis['total_profit'] = producer_analysis['avg_profit'] * producer_analysis['movie_count']

# Sort by score to find the best producers
best_producers = producer_analysis.sort_values(by='total_profit', ascending=False)

# Display the top producers
best_producers.head(10)
```